



Intro to Python and Pandas

DataFrames for Data Science

Sébastien Biass 

sebastien.biass@unige.ch

Earth Sciences

October 15, 2025

Background



We assume that you all followed Guy Simpson's Python crash course

pandas: A **package** for data manipulation and analysis handling **structured data**

- **Reading/writing data** from common formats (CSV, Excel, JSON, etc.)
- Handling **missing data**
- **Filtering, sorting, reshaping** and **grouping** data
- **Aggregating** data (sum, mean, count, etc.)
- **Time series support** (date ranges, frequency conversions)
- **Statistical operations**



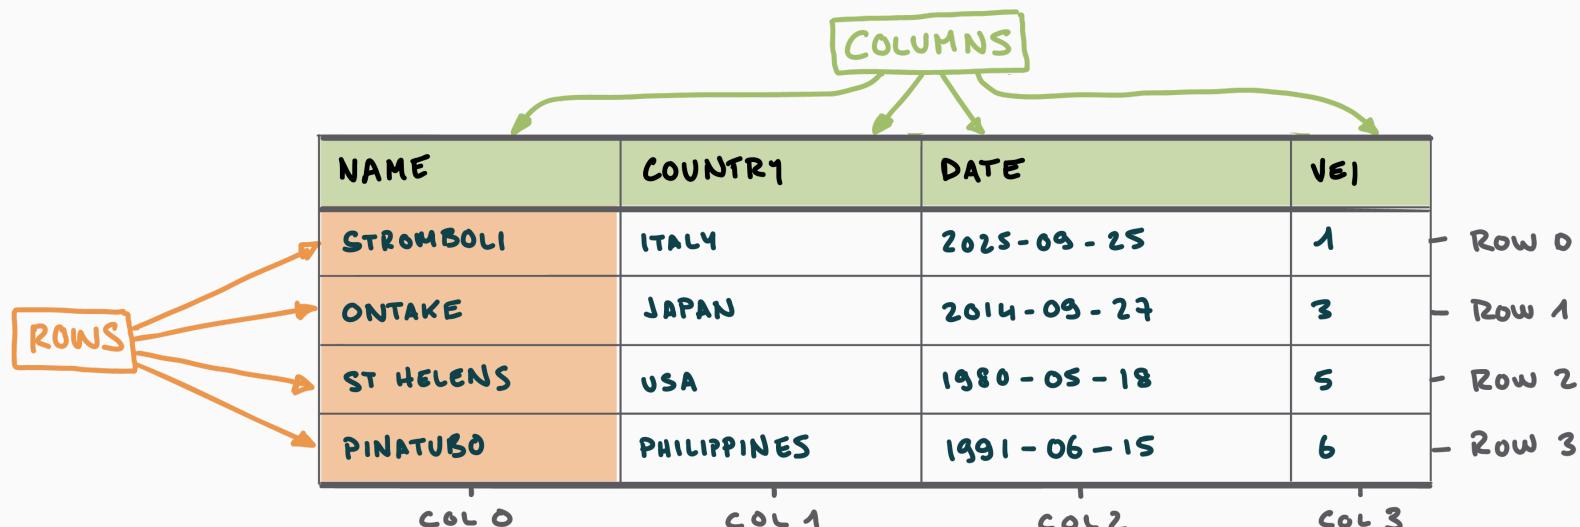
Today's objectives

Understand what is a **pandas** DataFrame and its basic anatomy

- How to load data in a DataFrame
- How to access data → *query by label/position*
- How to filter data → *comparison and logical operators*
- How to rearrange data → *sorting values*
- How to operate on data → *arithmetic and string operations*

Introduction to pandas

Anatomy of a DataFrame



- Similar to *Excel* → contains *tabular data* composed of **rows** and **columns**
- In *Excel*:
 - **Rows** are accessed using *numbers*
 - **Columns** are accessed using *letters*

Anatomy of a DataFrame

NAME	COUNTRY	DATE	VEI
STROMBOLI	ITALY	2025-09-25	1
ONTAKE	JAPAN	2014-09-27	3
ST HELENS	USA	1980-05-18	5
PINATUBO	PHILIPPINES	1991-06-15	6

INDEX = LABELS ALONG ROWS → df.index

COLUMNS = LABELS ALONG COLUMNS → df.columns

COL 0 COL 1 COL 2

- Unlike Excel, **rows** and **columns** can be labelled
 - Index refers to the label of the **rows**. In the index, **values are usually unique** - meaning that each entry has a different label.
 - Column refers to the label of - logically - the **columns**

Data structure

The dataset

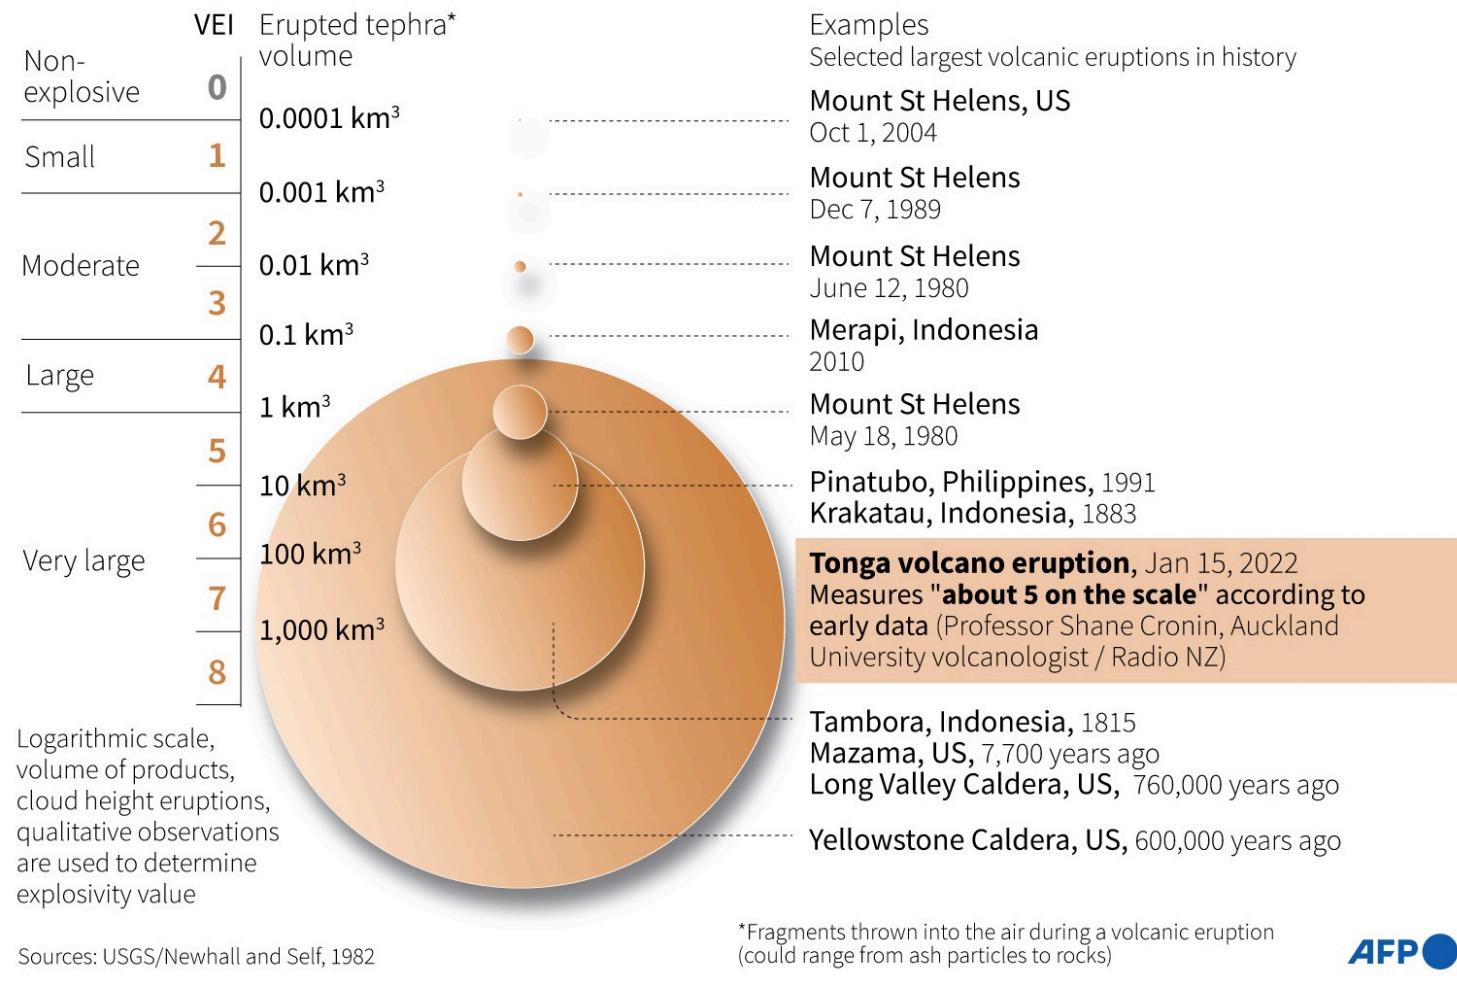
Synthetic dataset of **selected volcanic eruptions** → first 5 rows:

Name	Country	Date	VEI	Latitude	Longitude
St. Helens	USA	1980-05-18	5	46.1914	-122.196
Pinatubo	Philippines	1991-04-02	6	15.1501	120.347
El Chichón	Mexico	1982-03-28	5	17.3559	-93.2233
Galunggung	Indonesia	1982-04-05	4	-7.2567	108.077
Nevado del Ruiz	Colombia	1985-11-13	3	4.895	-75.322

Volcanic explosivity index (VEI)

Volcanic explosivity index

Measures the relative explosivity of volcanic eruptions





Setting up the notebook

- We start by importing the `pandas` library
- We import it under the name `pd` - which is faster to type!

```
1 # Import the required packages  
2 import pandas as pd
```



Setting up the notebook

- We then load the specified data with the `pd.read_csv()` function
- This returns a `DataFrame` object in a variable named `df`

```
1 # Import the required packages
2 import pandas as pd
3
4 # Read the data
5 df = pd.read_csv('https://raw.githubusercontent.com/e5k/Data-Science/main/Data/dummy_volcanoes')
```

Setting up the notebook

- We print some data for inspection with `df.head()`
- The functions are now directly called from the DataFrame `df` object

```
1 # Import the required packages
2 import pandas as pd
3
4 # Read the data
5 df = pd.read_csv('https://raw.githubusercontent.com/e5k/Data-Science/main/Data/dummy_volcanoes'
6
7 # Show the first 3 rows
8 df.head(3)
```

	Name	Country	Date	VEI	Latitude	Longitude
0	St. Helens	USA	1980-05-18	5	46.1914	-122.1956
1	Pinatubo	Philippines	1991-04-02	6	15.1501	120.3465
2	El Chichón	Mexico	1982-03-28	5	17.3559	-93.2233

Setting the index

NAME	COUNTRY	DATE	VEI
STROMBOLI	ITALY	2025-09-25	1
ONTAKE	JAPAN	2014-09-27	3
ST HELENS	USA	1980-05-18	5
PINATUBO	PHILIPPINES	1991-06-15	6

INDEX = LABELS ALONG ROWS → df.index

COLUMNS = LABELS ALONG COLUMNS → df.columns

COL 0 COL 1 COL 2

Row 0
Row 1
Row 2
Row 3

Setting the index

- ...for now, the index (\rightarrow **the first column**) is an *integer*
- This might be acceptable in datasets where the *label* is not important

```
1 # Show the first 3 rows
2 df.head(3)
```

	Name	Country	Date	VEI	Latitude	Longitude
0	St. Helens	USA	1980-05-18	5	46.1914	-122.1956
1	Pinatubo	Philippines	1991-04-02	6	15.1501	120.3465
2	El Chichón	Mexico	1982-03-28	5	17.3559	-93.2233

Setting the index

- Here we want to access the data using the **name of the volcano**
- We **set the index** using `set_index()`

```
1 # Set the index to the 'Name' column
2 df = df.set_index('Name')
3
4 # Show the first 3 rows
5 df.head(3)
```

	Country	Date	VEI	Latitude	Longitude
Name					
St. Helens	USA	1980-05-18	5	46.1914	-122.1956
Pinatubo	Philippines	1991-04-02	6	15.1501	120.3465
El Chichón	Mexico	1982-03-28	5	17.3559	-93.2233

Exploring data

- Here are some **basic functions** to review the structure of the dataset:

Function	Description
<code>df.head()</code>	Prints the <i>first</i> 5 rows of the DataFrame.
<code>df.tail()</code>	Prints the <i>last</i> 5 rows of the DataFrame.
<code>df.info()</code>	Displays some info about the DataFrame, including the number of rows (entries) and columns.
<code>df.shape</code>	Returns a list containing the number of rows and columns of the DataFrame.
<code>df.index</code>	Returns a list containing the index along the <i>rows</i> of the DataFrame.
<code>df.columns</code>	Returns a list containing the index along the <i>columns</i> of the DataFrame.



Functions vs attributes

- Functions** have parentheses → they **compute** something on `df`
- Attributes** do *not* have parentheses → they store some **parameter** related to `df`

Sorting data

- Sorting numerical, datetime or strings using `.sort_values`
- Importance of **documentation** to understand arguments

```
1 df.sort_values('VEI').head() # Sort volcanoes by VEI in ascending number
```

	Country	Date	VEI	Latitude	Longitude
Name					
Nyiragongo	DR Congo	2021-05-22	1	-1.5200	29.2500
Merapi	Indonesia	2023-12-03	2	-7.5407	110.4457
Ontake	Japan	2014-09-27	2	35.5149	137.4781
Kīlauea	USA	2018-05-03	2	19.4194	-155.2811
Etna	Italy	2021-03-16	2	37.7510	15.0044

Sorting data

- Sorting numerical, datetime or strings using `.sort_values`
- Importance of **documentation** to understand arguments

```
1 df.sort_values('VEI').head() # Sort volcanoes by VEI in ascending number  
2 df.sort_values('Date', ascending=False).head() # Sort volcanoes by eruption dates from recent
```

	Country	Date	VEI	Latitude	Longitude
Name					
Merapi	Indonesia	2023-12-03	2	-7.5407	110.4457
Cleveland	USA	2023-05-23	3	52.8250	-169.9444
Sinabung	Indonesia	2023-02-13	3	3.1719	98.3925
Nyiragongo	DR Congo	2021-05-22	1	-1.5200	29.2500
La Soufrière	Saint Vincent	2021-04-09	4	13.2833	-61.3875

Sorting data

- Sorting numerical, datetime or strings using `.sort_values`
- Importance of **documentation** to understand arguments

```
1 df.sort_values('VEI').head() # Sort volcanoes by VEI in ascending number
2 df.sort_values('Date', ascending=False).head() # Sort volcanoes by eruption dates from recent
3 df.sort_values('Country').head() # Also works on strings to sort alphabetically
```

	Country	Date	VEI	Latitude	Longitude
Name					
Calbuco	Chile	2015-04-22	4	-41.2972	-72.6097
Nevado del Ruiz	Colombia	1985-11-13	3	4.8950	-75.3220
Nyiragongo	DR Congo	2021-05-22	1	-1.5200	29.2500
Eyjafjallajökull	Iceland	2010-04-14	4	63.6333	-19.6111
Galunggung	Indonesia	1982-04-05	4	-7.2567	108.0771

Sorting data

- Sorting numerical, datetime or strings using `.sort_values`
- Importance of **documentation** to understand arguments

```
1 df.sort_values('VEI').head() # Sort volcanoes by VEI in ascending number
2 df.sort_values('Date', ascending=False).head() # Sort volcanoes by eruption dates from recent
3 df.sort_values('Country').head() # Also works on strings to sort alphabetically
4 df.sort_values(['Latitude', 'Longitude']).head() # Sorting using multiple columns
```

	Country	Date	VEI	Latitude	Longitude
Name					
Calbuco	Chile	2015-04-22	4	-41.2972	-72.6097
Agung	Indonesia	2017-11-21	3	-8.3422	115.5083
Merapi	Indonesia	2023-12-03	2	-7.5407	110.4457
Galunggung	Indonesia	1982-04-05	4	-7.2567	108.0771
Tavurvur	Papua New Guinea	2014-08-29	3	-4.3494	152.2847



Your turn!

- Go to <https://e5k.github.io/Data-Science/>
 - Class 1 > Data Structure

Class 1 > Data structure

Data structure

Let's get our hands dirty and start coding. Create a new Jupyter notebook following [this guide](#). You can copy fragments of the code, but make sure each code block is a different cell in your notebook. Also remember that you can add **Markdown** cells in between code cells, which are really useful to document your code.

The data we will use here is a `csv` file containing selected eruptions of the past 50 years. The first 5 rows of the data are illustrated in [Table 1](#).

Table 1: First 5 rows of the dataset.

Name	Country	Date	VEI	Latitude	Longitude
St. Helens	USA	1980-05-18 00:00:00	5	46.1914	-122.196
Pinatubo	Philippines	1991-04-02 00:00:00	6	15.1501	120.347

Querying data

Accessing data in a DataFrame

NAME	COUNTRY	DATE	VEI
STROMBOLI	ITALY	2025-09-25	1
ONTAKE	JAPAN	2014-09-27	3
ST HELENS	USA	1980-05-18	5
PINATUBO	PHILIPPINES	1991-06-15	6

INDEX = LABELS ALONG ROWS → df.index

COLUMNS = LABELS ALONG COLUMNS → df.columns

COL 0 COL 1 COL 2

Row 0
Row 1
Row 2
Row 3

Accessing data in a DataFrame

Option 1: label-based indexing

- Use the labels of **index** and **columns** to retrieve data
- Function to use: `df.loc`

`df.loc['ONTAKE', 'DATE']` → 2014-09-27

NAME	COUNTRY	DATE	VEI	
STROMBOLI	ITALY	2025-09-25	1	- Row 0
ONTAKE	JAPAN	2014-09-27	3	- Row 1
ST HELENS	USA	1980-05-18	5	- Row 2
PINATUBO	PHILIPPINES	1991-06-15	6	- Row 3

Accessing data in a DataFrame

Option 2: position-based indexing

- Use the positions of **index** and **columns** to retrieve data
- Function to use: `df.iloc`

`df.iloc [0,2] → USA`

NAME	COUNTRY	DATE	VEI	
STROMBOLI	ITALY	2025-09-25	1	- Row 0
ONTAKE	JAPAN	2014-09-27	3	- Row 1
ST HELENS	USA	1980-05-18	5	- Row 2
PINATUBO	PHILIPPINES	1991-06-15	6	- Row 3

COL 0 COL 1 COL 2

Label-based indexing: Rows

- Query a **row** with `.loc` → Use **square brackets** `[]`
 - Query the *row* for which the *index label* is `Calbuco`
 - Returns all *columns*

```
1 df.loc['Calbuco']
```

```
Country          Chile
Date    2015-04-22 00:00:00
VEI                  4
Latitude       -41.2972
Longitude      -72.6097
Name: Calbuco, dtype: object
```

→ Returns a `pd.Series`

```
1 df.loc[['Calbuco']]
```

	Country	Date	VEI	Latitude	Longitude
Name					
Calbuco	Chile	2015-04-22	4	-41.2972	-72.6097

→ Returns a `pd.DataFrame`

Label-based indexing: Rows

- Query **multiple rows** with `.loc`
 - Query the rows for which the *index labels* are `Calbuco` or `Taal`
 - Returns all *columns*

```
1 df.loc[['Calbuco', 'Taal']]
```

	Country	Date	VEI	Latitude	Longitude
Name					
Calbuco	Chile	2015-04-22	4	-41.2972	-72.6097
Taal	Philippines	2020-01-12	4	14.0020	120.9934

Label-based indexing: Columns

- **Query columns**

- Query the *columns* for which the *column labels* are `Country` or `VEI`
- Returns all rows

- **Option 1:** with `.loc`:

```
1 df.loc[:, ['Country', 'VEI']].head(3)
```

	Country	VEI
Name		
St. Helens	USA	5
Pinatubo	Philippines	6
El Chichón	Mexico	5

- **Option 2:** without `.loc`:

```
1 df[['Country', 'VEI']].head(3)
```

	Country	VEI
Name		
St. Helens	USA	5
Pinatubo	Philippines	6
El Chichón	Mexico	5

Label-based indexing: Rows and Columns

- Again, choice on whether to use `.loc` to query columns
- Option 1:** Columns are specified **inside** `.loc`:
- Option 2:** Columns are specified **outside** `.loc`:

```
1 df.loc[['Calbuco', 'Taal'], ['Country', 'V
```

```
1 df.loc[['Calbuco', 'Taal']][['Country', 'V
```

	Country	VEI
Name		
Calbuco	Chile	4
Taal	Philippines	4

	Country	VEI
Name		
Calbuco	Chile	4
Taal	Philippines	4

Position-based indexing

- Use **positions** instead of **labels**

`df.iloc [0,2] → USA`

NAME	COUNTRY	DATE	VEI	
STROMBOLI	ITALY	2025-09-25	1	Row 0
ONTAKE	JAPAN	2014-09-27	3	Row 1
ST HELENS	USA	1980-05-18	5	Row 2
PINATUBO	PHILIPPINES	1991-06-15	6	Row 3

COL 0 COL 1 COL 2

Position-based indexing: Rows

- Query a **row** with `.iloc`

→ Returns all *columns*

- One row** (first row):

```
1 df.iloc[0]
```

	Country	Date	VEI	Latitude	Longitude
Name					
St. Helens	USA	1980-05-18	5	46.1914	-122.1956

- Range of rows:** (rows 3-4):

```
1 df.iloc[2:4]
```

	Country	Date	VEI	Latitude	Longitude
Name					
El Chichón	Mexico	1982-03-28	5	17.3559	-93.223
Galunggung	Indonesia	1982-04-05	4	-7.2567	108.07

Position-based indexing: Rows

- Query rows from the end:
- Example:
 - Get the **last 5 rows** of the DataFrame:

```
1 df.iloc[-5:]
```

	Country	Date	VEI	Latitude	Longitude
Name					
La Soufrière	Saint Vincent	2021-04-09	4	13.2833	-61.3875
Calbuco	Chile	2015-04-22	4	-41.2972	-72.6097
St. Augustine	USA	2006-03-27	3	57.8819	-155.5611
Eyjafjallajökull	Iceland	2010-04-14	4	63.6333	-19.6111
Cleveland	USA	2023-05-23	3	52.8250	-169.9444

Position-based and label-based queries

- Mix position-based and label-based indexing:
 - **Rows** → *labels*
 - **Columns** → *positions*

```
1 df.iloc[0:5][['Country', 'VEI']]
```

	Country	VEI
Name		
St. Helens	USA	5
Pinatubo	Philippines	6
El Chichón	Mexico	5
Galunggung	Indonesia	4
Nevado del Ruiz	Colombia	3



Your turn!

- Go to <https://e5k.github.io/Data-Science/>
 - Class 1 > Querying data

Class 1 > Querying data

Querying data

Querying data from a DataFrame

Let's now review how we can access data contained in the DataFrame. This process, known as *indexing*, consists in specifying a row or a column (or ranges of rows and columns) where the data is stored. In `pandas`, there are two different ways to do that:

- By `label`: data is queried using the actual index/column name (e.g., the `VEI` column in the DataFrame above)
- By `location`: data is queried using the column location (e.g., the 3rd row)

Label-based indexing

`df.loc['ONTAKE', 'DATE']` → 2014-09-27

Filtering data

Boolean indexing

- **Filtering** data with boolean indexing
 - Returns either **True** or **False** depending on whether the **condition** is satisfied
- Example:

```
1 a = 1  
2 b = 2
```

Comparison operators:

Operator	Meaning	Example Result	
<code>==</code>	Equal to	<code>a == b</code>	<code>False</code>
<code>!=</code>	Not equal to	<code>a != b</code>	<code>True</code>
<code>></code>	Greater than	<code>a > b</code>	<code>False</code>
<code><</code>	Less than	<code>a < b</code>	<code>True</code>
<code>>=</code>	Greater than or equal	<code>a >= b</code>	<code>False</code>
<code><=</code>	Less than or equal	<code>a <= b</code>	<code>True</code>

Boolean indexing: Example

- Query all volcanoes where `VEI == 4`

```
1 df['VEI'] == 4
```

```
Name
St. Helens
False
Pinatubo
False
El Chichón
False
Galunggung
True
Nevado del Ruiz
False
Merapi
False
Ontake
False
Soufrière Hills
False
Etna
```

```
1 df.loc[df['VEI'] == 4]
```

	Country	Date	VEI	Latitude	Longitude
Name					
Galunggung	Indonesia	1982-04-05	4	-7.2567	108.0771
Taal	Philippines	2020-01-12	4	14.0020	120.9934
La Soufrière	Saint Vincent	2021-04-09	4	13.2833	-61.3875
Calbuco	Chile	2015-04-22	4	-41.2972	-72.6097
Eyjafjallajökull	Iceland	2010-04-14	4	63.6333	-19.6111

Boolean indexing: Strings

- **Filtering** also works with strings
 - Use **string comparison** operations

- **Example:**

```
1 df.loc[df['Country'] == 'Indonesia']
```

	Country	Date	VEI	Latitude	Longitude
Name					
Galunggung	Indonesia	1982-04-05	4	-7.2567	108.0771
Merapi	Indonesia	2023-12-03	2	-7.5407	110.4457
Agung	Indonesia	2017-11-21	3	-8.3422	115.5083
Sinabung	Indonesia	2023-02-13	3	3.1719	98.3925



Boolean indexing: Strings

- **Filtering** also works with strings
 - Use **string comparison** operations
- **String comparison operators:**

Operation	Example	Description
contains	<code>df['Name'].str.contains('Soufrière')</code>	Checks if each string contains a substring
startswith	<code>df['Name'].str.startswith('E')</code>	Checks if each string starts with a substring
endswith	<code>df['Name'].str.endswith('o')</code>	Checks if each string ends with a substring



Your turn!

- Go to <https://e5k.github.io/Data-Science/>
 - Class 1 > Filtering data

Class 1 > Filtering data

Filtering data

Now that we have reviewed how to access data, let's now see how to **filter** data using **boolean indexing**. For this, we need to review what are **comparison operators** ([Table 1](#)).

Comparison operators

Let's assume the following variables ([Listing 1](#)):

Listing 1: Variables used for illustrating logical operations

```
a = 1  
b = 2
```

Applying the comparison operators in [Table 1](#) will produce a variable of type `bool` -

Operations

Data management operations

- Common data management functions for pandas columns:

Operation	Example	Description
Round	<code>df['VEI'].round(1)</code>	Rounds values to the specified number of decimals
Floor	<code>df['VEI'].apply(np.floor)</code>	Rounds values down to the nearest integer
Ceil	<code>df['VEI'].apply(np.ceil)</code>	Rounds values up to the nearest integer
Absolute value	<code>df['VEI'].abs()</code>	Returns the absolute value of each element
Fill missing	<code>df['VEI'].fillna(0)</code>	Replaces missing values with a specified value

- Round the '**Latitude**' to two decimals → make sure you store the output!

```
1 df['Latitude'] = df['Latitude'].round(2)
```

Arithmetic operations

- Arithmetic operations on parts of the DataFrame (\rightarrow columns) using native Python arithmetic operators

Operation	Symbol	Example	Description
Addition	+	<code>df['VEI'] + 1</code>	Adds a value to each element
Subtraction	-	<code>df['VEI'] - 1</code>	Subtracts a value from each element
Multiplication	*	<code>df['VEI'] * 2</code>	Multiplies each element by a value
Division	/	<code>df['VEI'] / 2</code>	Divides each element by a value
Exponentiation	**	<code>df['VEI'] ** 2</code>	Raises each element to a power
Modulo	%	<code>df['VEI'] % 2</code>	Remainder after division for each element

- Divide `VEI` by 2 and store results a new column (\rightarrow `VEI_halved`)

```
1 df['VEI_halved'] = df['VEI'] / 2
```

Expanded arithmetic operations

- The range of arithmetic operations can be expanded using `numpy`

Operation	Symbol	Example	Description
Exponentiation	<code>np.power</code>	<code>np.power(df['VEI'], 2)</code>	Element-wise exponentiation
Square root	<code>np.sqrt</code>	<code>np.sqrt(df['VEI'])</code>	Element-wise square root
Logarithm (base e)	<code>np.log</code>	<code>np.log(df['VEI'])</code>	Element-wise natural logarithm
Logarithm (base 10)	<code>np.log10</code>	<code>np.log10(df['VEI'])</code>	Element-wise base-10 logarithm
Exponential	<code>np.exp</code>	<code>np.exp(df['VEI'])</code>	Element-wise exponential (e^x)

- `VEI` power 2 and store results a new column (\rightarrow `VEI_pow2`)

```
1 df['VEI_pow2'] = np.power(df['VEI'], 2)
```



Your turn!

- Go to <https://e5k.github.io/Data-Science/>
 - Class 1 > Operations

Class 1 > Operations

Operations

Data management operations

Pandas contains several helpful functions to manage and format numerical data ([Table 1](#)).

Table 1: Common data management functions for pandas columns.

Operation	Example	Description
Round	<code>df['VEI'].round(1)</code>	Rounds values to the specified number of decimals
Floor	<code>df['VEI'].apply(np.floor)</code>	Rounds values down to the nearest integer
Ceil	<code>df['VEI'].apply(np.ceil)</code>	Rounds values up to the nearest integer